

## Day 4 – PySpark Joins Deep Dive (L4-Level)

---

### ♦ 1. Sort-Merge Join (SMJ): When & Why

#### Concept:

- Used when both join sides are **large** and **already sorted** or **can be sorted** efficiently.
- Spark performs sort → merge → join.

#### When used:

- No broadcast possible (tables too big)
- Join key is sortable (e.g., string, int)
- Default fallback join type for large tables

#### Trade-offs:

- Sorting = CPU + shuffle-heavy
- Works well with **bucketing + sorting (explained in 7th question)** in Delta or Hive

#### Production Example:

- 100 GB orders joined with 200 GB transactions.
- SMJ used → sorted both by `order_id`.
- Sorting time optimized by **pre-sorting + bucketing** during write.



#### Quantified Improvement:

- Job duration: 82 min → 51 min
  - Shuffle read: 240 GB → 160 GB
  - GC reduced by 40% (fewer long-lived objects)
-

## ♦ 2. Shuffle Hash Join vs Broadcast Join

Join Type	When to Use	Trade-offs
<b>Broadcast</b>	One table is very small (<10MB)	Avoids shuffle entirely, fast. But OOM risk if size underestimated
<b>Shuffle Hash Join</b>	Mid-size joins, unsorted tables	Hashing and shuffling both sides; less CPU than SMJ, more memory pressure
<b>SMJ</b>	Both sides large, sorted	More stable, but CPU intensive

### Tuning Broadcast Join:

python

CopyEdit

```
spark.conf.set("spark.sql.autoBroadcastJoinThreshold", 10 * 1024 * 1024) # 10MB
```



### Quantified Example:

- Fact table (5GB) join with 20MB dimension table:
  - Shuffle join → 36 min
  - Broadcast join → 17 min (~2.1x faster)
  - Reduced shuffle read by 98%

---

## ♦ 3. Join Reordering & Optimizer Behavior

**Catalyst Optimizer** can **reorder joins** for performance using **cost-based stats**.

### What it does:

- Pushes filters early
- Joins smallest tables first
- May choose SMJ vs SHJ based on stats

### How to control:

python

CopyEdit

```
spark.conf.set("spark.sql.cbo.enabled", "true")
spark.conf.set("spark.sql.cbo.joinReorder.enabled", "true")
```



#### Impact:

- With CBO (Cost based Optimization) enabled:
    - DAG stages: 6 → 4
    - Job time: 38 min → 28 min
    - Filter pushdown saved 8GB scan
- 

## ♦ 4. Skewed Join Handling (Salting, Hints)

### Problem:

Join keys are skewed (e.g., 90% of rows have `user_id = 123`).

### Fixes:

#### 1. Salting:

- Add random suffix to skewed key
- Repartition both sides using salted key
- Post-join → remove salt

#### 2. Skew Join Hints (Spark 3.0+):

sql

CopyEdit

```
SELECT /*+ SKEW('key') */ ...
```

#### 3. Broadcast small side (if skewed side is big)



#### Real Improvement:

- Stage duration: 70 min → 26 min

- Skewed task retry count: 45 → 1
- Executor spill: 4.8 GB → 600 MB

## REPARTITION VS COALESCE IN SALTING

`coalesce()` is ideal for optimizing writes after filtering or aggregations when the data is small. In one job, we reduced output files from 200 to 5 using `coalesce(5)`, improving S3 read performance downstream and avoiding small file explosion (**i.e COALESCE is not helpful in salting**). But for salting or joins, `repartition()` remains the better choice due to shuffle-based distribution.

---

### ♦ 5. Null-safe Join (<=>)

#### Problem:

Default join behavior: `NULL != NULL`

#### Fix:

Use null-safe operator (<=>):

```
python
CopyEdit
df1.join(df2, df1.id <=> df2.id)
```

#### Use case:

- Auditing pipelines
- Data deduplication
- Slowly Changing Dimensions with nullable keys



#### Impact Example:

- Reduced mismatched rows from 8K → 0
- Improved match accuracy in dedup job by 100%

## Example 1: Deduplication with Nullable Keys



#### DataFrames:

python

CopyEdit

```
# df1
```

```
+-----+
```

```
| id |
```

```
+-----+
```

```
| 1 |
```

```
| 2 |
```

```
| NULL |
```

```
+-----+
```

```
# df2
```

```
+-----+
```

```
| id |
```

```
+-----+
```

```
| 1 |
```

```
| 3 |
```

```
| NULL |
```

```
+-----+
```

**✗ Using default join (==):**

python

CopyEdit

```
df1.join(df2, df1.id == df2.id, "inner").show()
```

## Result:

diff

CopyEdit

```
+----+----+
```

```
|id |id |
```

```
+----+----+
```

```
| 1 | 1 |
```

```
+----+----+
```

- `NULL == NULL` is **false**, so they don't match.

---

## ◆ 6. Broadcast Join Threshold Tuning

### Config:

python

CopyEdit

```
spark.sql.autoBroadcastJoinThreshold = 10MB # default
```

### When to increase:

- You know the dimension table is 15–20MB and stable
- You have enough executor memory (~8GB per executor)

### Caution:

Setting too high = OOM risk



### Example:

- Raising threshold from 10MB → 20MB enabled broadcast join
- Job time: 33 min → 16 min

- Reduced stage count: 5 → 2
- 

## ♦ 7. Bucketing for Joins

### Concept:

- Pre-shuffle data into same number of buckets on same key
- Join doesn't need to shuffle again → saves I/O

### Steps:

python

CopyEdit

```
df.write.bucketBy(8, "user_id").sortBy("user_id").saveAsTable(...)
```

### Requirements:

- Bucketing enabled in session
- Same bucket count + key on both sides



### Example:

- 300 GB → 300 GB join:
  - Without bucketing: 75 min
  - With bucketing: 39 min (~48% faster)
  - Shuffle read reduced by ~60%

**? If I save both tables with `bucketBy()` and `saveAsTable()`, and then I do a join using PySpark DataFrame API (not SQL), will shuffle happen?**

---

✗ **Short Answer: YES — Spark will still perform a shuffle.**

---

## Why?

Because when you do:

python

CopyEdit

```
df1 = spark.table("bucketed_orders")  
  
df2 = spark.table("bucketed_payments")  
  
df1.join(df2, "order_id").explain()
```

- ➡ Even though the tables are saved with buckets,
  - ➡ The **PySpark DataFrame join API does not use bucketing metadata** from the metastore.
  - ➡ So Spark **treats it like a normal join → shuffles both sides.**
- 

## **Bucketing optimization is triggered only when using SQL-based queries**

sql

CopyEdit

```
SELECT * FROM bucketed_orders o  
  
JOIN bucketed_payments p  
  
ON o.order_id = p.order_id
```

➡ In this case, **Catalyst planner checks bucketing info from Hive metastore**, verifies that both sides:

- are bucketed on same column



- have same number of buckets
- are sorted

➡ Then it can **skip shuffle** in Sort-Merge Join

## Summary Table

Join Method	Bucketing Used?	Shuffle Avoided?
<code>spark.sql("SELECT * FROM ...")</code>	✅ Yes	✅ Yes
<code>df1.join(df2, "key")</code>	❌ No	❌ No
<code>df1.join(df2, "key").explain()</code>	shows Exchange	shows shuffle

## Final Takeaway for Interview:

"Even if I save tables using `.bucketBy()` and `.saveAsTable()`, Spark will only avoid shuffle in joins if I query using `spark.sql()`. The PySpark DataFrame API does **not leverage bucketing metadata**, so it still performs a shuffle."

## ! Is bucketing useful in PySpark if you don't use `spark.sql()`?

❌ No — bucketing is practically useless in plain PySpark DataFrame joins.

## Why?

Because:

- PySpark's `.join()` **does not read or use** bucketing metadata.
  - Even if data was saved using `.bucketBy()` and looks physically bucketed (`bucket_00000`, etc.), **Spark treats it like regular data** unless:
    - You use `.saveAsTable()`
    - AND query using `spark.sql()`
- 

## So When Is Bucketing Useful?

Scenario	Bucketing Works?	Shuffle Avoided?
PySpark <code>.join()</code> on bucketed files	✗ No	✗ No
PySpark <code>.join()</code> on <code>spark.table()</code>	✗ No	✗ No
<code>spark.sql()</code> on saved bucket tables	✓ Yes	✓ Yes

---

## When Should You Use Bucketing in PySpark?

Only when:

- You're writing data to Hive or Delta Lake tables

- You plan to query them via SQL (not DataFrame API)
  - You want to optimize **Sort-Merge Joins** on large datasets
  - You use tools like **Databricks** or **Hive-aware** environments
- 

## Otherwise?

"For standard PySpark DataFrame code (non-SQL), **bucketing is just an expensive write operation that gives you no runtime benefit.**"

---

## Practical Alternatives to Bucketing in PySpark:

If you're working entirely in PySpark, and want join optimizations:

1. Use **Broadcast Joins** (for small dimension tables)
2. **Repartition** on join key before join:

```
python
CopyEdit
df1 = df1.repartition("user_id")

df2 = df2.repartition("user_id")

joined = df1.join(df2, "user_id")
```

3. Use **co-partitioned** writes if you're persisting intermediate stages
- 

## ♦ 8. Join Performance Troubleshooting

What to check in Spark UI:

- High shuffle read/write = fallback to SMJ or SHJ
- High GC = executor too large

- Skewed task duration = data skew
- Stage retries = OOM or straggler node

#### Fixes:

- Use `.hint("broadcast")` or `.hint("merge")`
- Repartition based on join key
- Avoid joining wide tables with UDFs on join key

#### Fix Story:

- Added `.hint("broadcast")` to stable 8MB table
- Reduced task count from 12K → 2.1K
- Job time: 42 min → 19 min

## 9. Problems with Excessive `repartition(n)`

### 1. GC Stalls

- Too many partitions → too many **tasks** → **too many objects** → **more frequent garbage collection**
  - Executors may start **spending 30–50% time in GC** instead of real work
  - Especially bad with large executor memory (heap bloat = longer GC pause)
- 

### 2. Scheduler & Task Overhead

- Spark's driver must track each task — too many = **driver memory pressure**
  - Too many small tasks = overhead in task scheduling, context switching
- 

### 3. Small File Problem

- If you `repartition(1000)` and then write, you'll get **1000 files** — hurts read performance
  - Common S3/ADLS problem: too many tiny files → listing + read latency explodes
- 

## Best Practice for Repartitioning:

Guideline	Value
Partitions = 2x–4x total cores	e.g., 200 cores → ~400–800 partitions
Use <code>spark.default.parallelism</code>	As a base value
Avoid > 10,000 partitions	Unless truly needed for scale