

Q1. What is a DAG in Spark? How is it different from a physical execution plan?

1. Problem Context:

Spark represents your entire job as a **Directed Acyclic Graph** (DAG) — a logical flow of transformations.

2. Concept Explanation:

- **Logical DAG:** Formed from transformations before an action. It's like:
`df.filter().groupBy().map()`
- **Physical Plan:** Once an action triggers, Spark converts the DAG into **Stages** and **Tasks**.
- **Stage:** A set of pipelined operations (e.g., map/filter) that don't require a shuffle
- **Task:** Unit of work per partition in a stage

3. Design & Trade-offs:

- DAG is **recomputed from lineage** if there's a failure
- If you persist intermediate results, Spark will not recompute from scratch
- Checkpointing truncates the lineage when it's too long or risky (e.g., >100 stages deep) (9th Question explained deeply)

4. Real-world Example:

"In one pipeline, lineage caused 4 re-computations of expensive joins when a worker failed. We persisted before the join and saved 40 minutes per failure."

5. Takeaway:

DAG helps with fault tolerance, but deep lineage without persistence = risk of recomputing huge pipelines.

Q2. How does Spark handle memory, GC, and spill to disk?

1. Problem Context:

Long-running jobs with wide transformations can spill intermediate data to disk or trigger GC stalls.

2. Concept Explanation:

- **Spark memory = Storage + Execution**
 - `spark.memory.fraction`: total usable memory (default 60%)
 - `spark.memory.storageFraction`: % reserved for cache vs computation
- **GC tuning**: Triggered when memory is tight — causes task delays
- **Spill to disk**: Happens during wide ops (joins, aggregations) when memory isn't enough

3. Design & Trade-offs:

- Use Kryo serialization for faster (and smaller) object encoding
- Monitor GC time in Spark UI → If high, reduce memory per executor or use more cores
- Don't oversubscribe memory — better to **increase executors** with lower memory than a few with high memory

4. Real-world Example:

“One job took 5h due to GC stalls. GC time was 50%. Switched from Java to Kryo, reduced executor memory to 4g, added more executors → GC time dropped to 3%, job time to 40 min.”

5. Takeaway:

GC overhead = hidden killer. Spark UI → Executors tab is your debugging friend.

Q3. How does shuffle work and what causes shuffle spills?

1. Problem Context:

Wide transformations (join, groupBy) require shuffling — expensive, disk-heavy, and often bottlenecks.

2. Concept Explanation:

- Data is **written by mappers** → **read by reducers**, creating temporary files
- **Spill occurs** when memory is insufficient → Spark spills sort buffers to disk
- Shuffle files are cleaned up *only after* job completes (except with external shuffle service)

3. Design & Trade-offs:

- Tune with `spark.sql.shuffle.partitions` (default 200 is often too low)
- Enable `spark.shuffle.compress=true` to save I/O
- External shuffle service keeps shuffle data across executor loss

4. Real-world Example:

“At Uber-scale job, shuffle read = 1.5 TB, spilled data = 700 GB. Increased partitions from 200 to 800, added compression → spill dropped to 180 GB, job time cut by 2h.”

5. Takeaway:

Shuffle = most expensive phase. Minimize it or optimize it — or your job dies at 95% progress.

Q4. A stage is stuck at 99% for 15 mins. What do you do?

1. Initial Suspects:

- A task is **skewed** → taking longer than others
- GC overhead or memory pressure on that executor
- Shuffle fetch failure or disk spill on slow node

2. Actions Taken:

- Check Spark UI → see stage DAG, task duration, skewed input size
- Look for tasks with 10x runtime or shuffle read
- Use `spark.sql.adaptive.skewJoin.enabled=true` in Spark 3+
- Repartition by range to reduce skew

3. If reproducible:

- Add salting or bucket join strategy
- Use persistent caching for upstream if re-computations are observed

4. Real-world Fix:

“One job was stuck on task 17/100 for 25 mins. Task was reading 20x more data.”

Salting + AQE skew join enabled solved it — saved 70% runtime.”

5. **Takeaway:**

Always check Spark UI → task duration, input size, GC time, shuffle reads.

✅ **Q5. What is the External Shuffle Service (ESS), and when is it required in production?**

1. **Problem Context:**

When an executor crashes or gets decommissioned mid-job, its shuffle files can be lost — causing job failure or retries.

2. **Concept Explanation:**

- ESS is a daemon that runs outside the executor and manages shuffle file lifecycle
- Allows **other executors** or **speculative tasks** to fetch shuffle data even if original executor dies
- **Required when:** Using dynamic allocation (executor churn is common), or in long shuffle-heavy jobs

3. **Design & Trade-offs:**

- ESS increases fault tolerance
- Needs correct config (`spark.shuffle.service.enabled=true`, port setup)
- Extra daemon process = more ops overhead

4. **Real-world Example:**

“A job crashed at stage 6/7 after 3 hours — root cause: lost executor with intermediate shuffle files. Enabled ESS → retried only failed tasks, not whole job → saved 2h runtime.”

5. **Takeaway:**

If you're running on YARN or with dynamic allocation — ESS is non-negotiable for resilience.

✓ Q6. What is Task Speculation in Spark? When do you use or avoid it?

1. Problem Context:

In large clusters, a few straggling tasks can delay job completion significantly.

What is Task Speculation in Spark?

Speculative execution is Spark's way of saying:

"Some tasks are running too slow. Maybe they're on a bad node or got unlucky. Let's start a duplicate of that task elsewhere and keep whichever finishes first."

2. Concept Explanation:

- Speculation means **launching duplicate tasks** for slow-running ones
- Whichever completes first, result is accepted
- Enabled via `spark.speculation=true`

3. Design & Trade-offs:

- ✓ Good when slow nodes cause tail latency
- ✗ Bad if slowness is due to external API, DB, or skew (wastes resources)
- Causes extra compute pressure if misused

4. Real-world Example:

"Job had 99/100 tasks complete in 5 mins; last one took 27 min. Enabled speculation → duplicated last task → job time reduced to 6 mins. Later avoided on DB-write job where speculation caused duplicate inserts."

5. Takeaway:

Use it **only for CPU-bound workloads** with large task parallelism and occasional tail latency.

Task speculation = Spark **duplicates slow-running tasks** (called *stragglers*) and **launches them on different executors**, not on the same one.

! Key Point:

Speculative tasks are always run on different executors, not on the same one.

Why?

- If the original executor is slow (e.g., GC issue, disk I/O, data skew), launching the same task again **on the same executor** would **not help**.
- Spark instead tries to launch the duplicate **on a faster, healthier executor** that still has available CPU cores and memory.

✓ Q7. GC Overhead Limit Exceeded – What Causes It? How Do You Fix It?

1. Problem Context:

Spark job fails due to JVM error: **GC Overhead Limit Exceeded** — meaning >98% of CPU time is spent in GC, and <2% heap is recovered.

2. Causes:

- Too little memory per executor
- Too many objects (especially with map-side aggregations or large joins)
- Poor serialization (e.g., Java serializer with complex nested data)

3. Fixes:

- Switch to **Kryo serializer**
(`spark.serializer=org.apache.spark.serializer.KryoSerializer`)
- Reduce memory per executor to allow more containers
- Tune GC: G1GC for large heaps
- Repartition to reduce data per task

4. Real-world Fix:

“1.2B-row join caused GC spike. GC logs showed >90% time in GC. Switched to Kryo, cut executor memory from 12g → 6g, increased executors → job time dropped from 3.5h to 48 mins.”

5. Takeaway:

GC tuning isn't optional at scale. Use Spark UI → GC Time per executor, or review

GC logs.

✓ Q8. How Do You Decide Spark Executor Config in Production?

1. Problem Context:

Wrong executor settings = wasted resources, OOMs, slow jobs.

2. Rules of Thumb:

- Start with: `spark.executor.instances`, `spark.executor.memory`, `spark.executor.cores`
- Use **1 executor per node** for large-memory nodes, or **2–5 per node** on standard
- Each executor → 3–5 cores ideally, too many = parallelism issues
- Memory = ~90% of available *container* memory – 1 GB for overhead

3. Real-world Config Comparison:

Config	Result
1 exec, 16g memory, 8 cores	✗ GC pause spikes, 1 task slow
4 execs, 4g memory, 2 cores	✓ Better GC, parallelism improved
Kryo + 8g + 3 cores + 4 execs	✓ Final config — balanced

4.

Tuning Tip:

Add `--conf spark.memory.fraction=0.7` if storage vs compute is unbalanced.

5. Takeaway:

Every workload is different. Benchmark, test in staging, monitor GC, shuffle, CPU.

9) CHECKPOINTING

Final State:

Aspect	Location	Notes
✓ New lineage	From HDFS file	Starts from checkpointed file as the base
✗ Old lineage (1–49)	Discarded (from DAG)	Not used anymore after checkpoint
✓ Checkpoint data	Stored in HDFS	Parquet-like format, managed by Spark

What does "new lineage" mean?

- Spark now **treats the checkpointed file as a fresh input source** (like reading from Parquet)
- All further transformations (51–100) are applied **on this file**, not on the original raw source
- So if failure happens later (e.g. transformation 70 fails), Spark **does NOT go back to step 1**
- It starts from the **checkpointed HDFS file**

ONLY AFTER AN ACTION CHECKPOINT SHOULD BE ENABLED, NOT BEFORE THAT