

♦ **Q1. Compare Parquet, Avro, and ORC in terms of structure, performance, and use**

Feature / Format	Parquet	Avro	ORC
Data Layout	Columnar	Row-based	Columnar
Compression	Good	Moderate	Best-in-class
Predicate Pushdown	✅ Yes	❌ No	✅ Yes
Splittable	✅ Yes	✅ Yes	✅ Yes
Schema Evolution	✅ Supported, but limited flexibility	✅ Best-in-class	✅ Supported, less flexible than Avro
Read Efficiency	✅ Great for column-based queries	❌ Slower for column queries	✅ Excellent for Hive workloads
Write Efficiency	Moderate	✅ Fast	Moderate
Support in Spark	✅ Native support	✅ Native support	✅ Supported (but better in Hive setups)
Best Use Cases	Analytics, Data Lakes, Spark SQL	Kafka pipelines, CDC, serialization	Hive warehouses, Hadoop-native systems

✅ **Why Avro (Row-based) Doesn't Support Predicate Pushdown:**

- **Avro is a row-based format**, meaning it stores entire rows together (all columns of one record).
- To filter on a specific column (e.g., **WHERE age > 30**), Spark must **read the full row** first — it cannot skip reading other columns efficiently.

🔍 **Predicate Pushdown Needs Columnar Layout**

Formats like **Parquet** and **ORC** are **columnar**, which means:

- Each column is stored separately.
- So Spark can **read only the needed column** (e.g., **age**) and **skip the rest**.
- This enables **predicate pushdown**: filters are pushed down to the storage layer → less I/O → faster queries.

IMP : SPARK STORES DATA AS AN CATALYST DATAFRAME, NOT AS AVRO/ PARQUET

Stage	Format Used	Notes
Kafka → Spark	Avro → Catalyst DF	Avro decoded into Catalyst memory
Computation in Spark	Catalyst (RAM)	No Parquet here
Writing to Delta Lake / S3 / HDFS	Parquet	Controlled by sink format

What	Format
Input from Kafka	Avro (or JSON, CSV etc.)
Spark computation	Catalyst DF (in-memory)
Spill to disk	Internal binary (not Parquet/Avro)
Sink (e.g., Delta)	Parquet + _delta_log

Stage	Stored as	Push-down possible?
Kafka Avro → Spark (ingestion)	Avro bytes → Catalyst rows	No
Parquet → Spark (read)	Parquet → Catalyst rows	Yes, before rows hit memory

Predicate push-down is a *read-time* optimisation.

It works only if the source format advertises and supports it (Parquet/ORC/JDBC).

Once data is inside the Catalyst engine , Spark just filters in RAM—format no longer matters. **(i.e if it is Parquet - Catalyst DF knows it supports predictive pushdown, for further join/, etc operations , PP is supported, else Not)**

Spark knows during read time if pushdown is possible.(i.e from avro/Parquet/CSV/JSON just before getting converted to catalyst DF)

If source = Avro → pushdown is not possible → everything comes into memory.

-
- ♦ Q2. What is predicate pushdown? Which formats support it?

✓ Concept:

- Lets Spark **skip unnecessary blocks** of data at the file reader level
- **Parquet/ORC** support it due to columnar structure
- **Avro, JSON, CSV** — no real predicate pushdown

🧠 Example:

python

CopyEdit

```
df.filter("region = 'US'")
```

Only the 'region' column is read from disk in **Parquet**.

♦ Q3. What is the small file problem in Spark? How do you fix it?

Issue	Impact
Too many small files (e.g., 100k files <1MB)	File lookup overhead, task explosion, poor parallelism

✅ Fixes:

- Write fewer, larger files using `.repartition(n)` (When shuffling) or `.coalesce(n)` (no shuffling)
- Use `spark.sql.files.maxPartitionBytes` config
- Use Delta's `OPTIMIZE` or compaction jobs

🧠 Real Scenario:

Your streaming job writes 1 file per trigger → leads to 10,000+ small files on S3. Fix by batching and compacting downstream.

♦ Q4. Explain **inferSchema**, **mergeSchema**, and their performance impact.

Option	Role	Impact
inferSchema	Guess column types from data	Costly for large datasets
mergeSchema	Combine schemas across files (Parquet)	Useful for evolution, but very expensive on large datasets

✅ **Best Practice:**

- Always define schemas explicitly in production
- Avoid **mergeSchema = true** in mainline ingestion jobs

🔍 Explanation per file format:

File Format	Schema Inference Needed?	inferSchema Used?	Notes	📄
CSV	✅ Yes (no schema stored)	✅ Yes (optional)	Must infer or define manually	
JSON	✅ Yes	✅ Yes (optional)	Can infer, but expensive	
Avro	❌ No (schema stored in header)	❌ Not needed	Schema is embedded in file	
Parquet	❌ No (schema stored in footer)	❌ Ignored	Schema is read directly	
ORC	❌ No	❌ Not needed	Like Parquet, optimized for analytics	

Avro/ Parquet/ ORC - **file formats** that stores schema **internally** in the file footer.


♦ Q5. How do you handle corrupted/bad records in ingestion?

♦ **What Are Bad Records?**

Rows that:

- Have **incomplete data** (e.g., missing values for required columns)
- Have **type mismatches** (e.g., string in an integer column)
- Are **malformed** (e.g., bad JSON format)

✅ Available Options in `.option("mode", ...)`

Mode	Behavior	
"PERMISSIVE"	(Default) Tries to parse rows. If bad, puts <code>null</code> in corrupt columns and logs in <code>_corrupt_record</code> .	
"DROPMALFORMED"	Silently drops bad rows . No trace.	
"FAILFAST"	Fails immediately on first corrupt record.	

📌 Example: Reading Malformed JSON

File: data.json

json

Copy Edit

```
{ "id": 1, "name": "Alice" }
{ "id": 2, "name": "Bob" }
{ "id": 3, "name": "Charlie"
{ "id": 4, "name": "David" }
```

Row 3 is missing a closing brace.

Code:

python

Copy Edit

```
df = spark.read \
    .option("mode", "PERMISSIVE") \
    .option("columnNameOfCorruptRecord", "_corrupt_record") \
    .json("data.json")

df.show(truncate=False)
```

Output:

pgsql

Copy Edit

```
+---+-----+-----+
| id | name | _corrupt_record |
+---+-----+-----+
| 1  | Alice | null            |
| 2  | Bob   | null            |
| null | null | {"id": 3, "name": "Charlie" |
| 4  | David | null            |
+---+-----+-----+
```



🔑 Takeaway:

- **Never silently drop** malformed data in production.
- Always **log, isolate, or quarantine** bad records.
- Use **"PERMISSIVE"** with **_corrupt_record** and route them to **Delta Lake quarantine tables**.

♦ **Q6. What are common ingestion pipeline optimizations?**

Area	Optimization
Partitioning	Use <code>partitionBy("dt")</code> when writing
File size	Aim for 100–500 MB per file
Schema	Provide <code>.schema()</code> instead of inferring
Input splits	Avoid deep nested folders unless needed
Compression	Use <code>snappy</code> (default for Parquet) for fast compression/decompression

♦ **Q7. What's the difference between `.save()`, `.saveAsTable()`, and `.insertInto()`?**


Method	Use When	Behavior
<code>.save()</code>	Write to path	Just writes files
<code>.saveAsTable()</code>	Register DataFrame as Hive table	Creates table + writes data

`.insertInto()` Insert into existing table

Table must exist, schema must match exactly


♦ Q8. What happens internally when Spark reads a file?

♦ Step-by-Step Internal Workflow

Step	What Happens	Example	
1	Spark contacts the file system (like S3, HDFS, DBFS)	When you do <code>spark.read.parquet("s3://bucket/path")</code> , Spark queries S3 to get metadata (file list, sizes, locations).	
2	File is split into logical partitions	A 1 GB Parquet file might be split into 8 partitions (~128 MB each by default) based on <code>spark.sql.files.maxPartitionBytes</code> .	
3	Schema is read (or inferred)	- For Parquet: Schema is read from the file footer (fast). - For CSV/JSON: Schema might be inferred by reading sample rows (slow).	
4	Executors read file blocks	Each partition is sent to a task, which runs on an executor. That task reads the corresponding file split (block) .	
5	Deserialize → Catalyst DataFrame	Data is deserialized into Spark's internal Catalyst format (binary, optimized).	
6	Lazy evaluation	No actual reading/transformation happens yet — Spark just builds the logical plan. Actual work happens only when you call an action like <code>.show()</code> , <code>.collect()</code> , <code>.write()</code> .	

♦ Q9. What's a good ingestion-to-storage pattern?

✅ Stage-wise Explanation

Layer	Format Used	What It Stores	Why It's Needed	
Bronze	Avro	Raw Kafka data (possibly nested, unclean)	Keeps original source; Avro supports schema evolution well	
Silver	Parquet	Flattened, cleaned, enriched data	Columnar format → fast analytical reads, partitioning works well	
Gold	Delta	Aggregated, optimized data for consumption	Supports time travel , MERGE , exactly-once , BI dashboards	